

# Clinical Documentation Assistant

---

Privacy-Focused AI Application

Built on MedGemma (HAI-DEF)

**Hackathon Team — Healthcare AI Track**

January 18, 2026

## 1 Introduction & Context

### 1.1 The Privacy Imperative in Healthcare AI

Patient data is the most sensitive digital asset. Cloud-based AI requires data to leave institutional boundaries, creating unacceptable compliance gaps and trust erosion. Our on-premise approach eliminates these risks entirely.

### 1.2 Why MedGemma?

- **Medical Domain Mastery:** Trained on PubMed, MIMIC-IV, clinical notes
- **Open Weights:** Deploy anywhere—no vendor lock-in
- **Efficient Variants:** 2B/7B models for clinic hardware
- **Research-Backed:** HAI-DEF family from reputable medical AI researchers

#### Key Metrics & Differentiators

- **Privacy:** 100% on-premise processing
- **Speed:** 60% faster documentation
- **Cost:** 90% cheaper than cloud subscriptions
- **Accuracy:** MedGemma fine-tuned for clinical language
- **Compliance:** HIPAA/GDPR ready by design
- **Offline:** Zero internet dependency

## 2 Product Definition

#### Product Vision

To empower every clinician with an AI assistant that **thinks like a medical expert** but **operates like a trusted local tool**—eliminating documentation burden while guaranteeing patient data never leaves the room.

### 2.1 Value Proposition Matrix

For Clinicians	For Institutions	For Patients
60% less time on notes	No cloud fees or BAAs	Data never shared externally
Reduced cognitive load	Full compliance control	More clinician eye contact
Better note consistency	Predictable ownership costs	Trust in privacy protection

## 3 Problem Statement

### 3.1 The Documentation Crisis

- **2:1 Ratio:** 2 hours charting per 1 hour patient care
- **70% Burnout:** Linked to administrative burden
- **\$4.6B/yr:** Cost of physician burnout in US
- **15+ Clicks:** Average per EHR note entry

### 3.2 Why Existing Options Fail

- **Cloud AI Scribes:** Privacy violations, latency, recurring costs
- **EHR Templates:** Inflexible, promote note bloat

- **Human Scribes:** \$30k+/year, variable quality
- **Dictation Software:** No medical intelligence

## 4 Target Users & Personas

### Primary Persona: Dr. Avery Chen

*Internal Medicine Physician, Community Clinic* Pain Points:

- Sees 25 patients/day, spends evenings on documentation
- Worried about cloud AI data sharing with tech companies
- Clinic IT budget limited, can't afford enterprise solutions

Needs:

- Fast, accurate SOAP notes during/after visits
- Absolute data privacy guarantee
- Simple interface, minimal training

## 5 Use Cases & Workflows

### 5.1 Core Workflow: SOAP Note Generation

#### Patient Visit → Final Note in 4 Steps

**Step 1: Input:** Clinician types/dictates raw observations

**Step 2: AI Processing:** MedGemma structures data locally

**Step 3: Draft Review:** Edit AI-generated SOAP note

**Step 4: Finalize:** Sign and export to EHR (manual)

### 5.2 Use Case Spectrum

Use Case	Input	Output
SOAP Notes	Free-text symptoms, exam findings	Structured SOAP format
Progress Notes	Bullet points from visit	Coherent narrative summary
Referral Letters	Patient data + reason	Professional letter draft
Discharge Summaries	Hospital course highlights	Structured discharge note

## 6 Functional Scope

### 6.1 Core Features

- ✓ Secure local web interface for clinician input (text only)
- ✓ On-premise MedGemma deployment (2B/7B parameter models)
- ✓ SOAP, progress note, and summary generation from free text
- ✓ Full editing capabilities—AI only provides drafts
- ✓ 100% offline operation—no internet required
- ✓ Plain text export for EHR manual import

### 6.2 Explicit Exclusions

- ✗ Medical diagnosis or decision support
- ✗ Direct audio processing (separate privacy concern)
- ✗ Automated EHR integration (institution-specific)

- ✗ Patient identification or demographics management
- ✗ Persistent storage without explicit consent

### 6.3 Ethical Guardrails


- **Human-in-the-Loop:** Clinician responsible for all final content
- **Transparency:** Clear "AI Draft" watermarks on all generated text
- **Non-Diagnostic:** System cannot suggest diagnoses or treatments
- **Data Minimization:** By default, no persistence of patient data

## 7 AI Design & Architecture

### AI Architecture

**Core Principle:** Privacy through architecture  
**Model:** MedGemma-2B (optimized for clinic hardware)  
**Deployment:** Docker container on local server/workstation  
**Processing:** All inference on-premise, zero external calls  
**Security:** Full disk encryption, user authentication layer

### 7.1 System Architecture



figureSystem Architecture – Privacy-First On-Premise Deployment

## 8 AI Interaction Design

## 8.1 Input/Output Schema

Input (Clinician)	→	Output (MedGemma)
”72F, diabetic, foot ulcer 2wk, no fever”	→	SUBJECTIVE: 72yo F with DM2 presents with 2-week history...
”BP 150/90, HR 88, afebrile”	→	OBJECTIVE: Vitals: BP 150/90, HR 88, T 98.6F...
”Stable chronic wound, no signs of systemic infection”	→	ASSESSMENT: Chronic diabetic foot ulcer, currently stable without evidence of active infection or osteomyelitis.
”Continue insulin, add ABX”	→	PLAN: 1. Continue home insulin regimen...

## 8.2 Safety Controls

- **Temperature=0.1:** Minimizes randomness, maximizes consistency
- **Input Sanitization:** Remove PHI patterns before processing
- **Max Tokens:** Limit output length to prevent verbosity
- **Keyword Filters:** Flag unusual content for review

## 9 Prompt Structure

### 9.1 Draft Prompt Design

The prompts for MedGemma are designed to be concise, structured, and focused on documentation tasks. A typical prompt follows this template:

You are a clinical documentation assistant. Structure the following raw notes into a standard S

Raw notes: [INSERT CLINICIAN INPUT HERE]

Output format:

SUBJECTIVE: [Summary of patient complaints]

OBJECTIVE: [Exam findings and vitals]

ASSESSMENT: [Neutral restatement of facts]

PLAN: [Organized action items]

This structure ensures the AI processes input safely and outputs only organized content.

## 10 Data & Evaluation

### 10.1 Example Inputs/Outputs

Here are sample interactions to demonstrate MedGemma's performance:

Input	→	Output
”45M, chest pain since morning, no shortness of breath, EKG normal”	→	SUBJECTIVE: 45-year-old male reports chest pain starting this morning... OBJECTIVE: EKG within normal limits...
”Follow-up for hypertension, meds compliant, BP 130/80”	→	SUBJECTIVE: Patient compliant with medications... OBJECTIVE: BP 130/80...

## 10.2 Sample Medical Report

### Sample SOAP Note Output

**SUBJECTIVE:** 65-year-old female with history of hypertension presents with fatigue and mild headache for 3 days. No chest pain or dizziness.

**OBJECTIVE:** BP 145/85, HR 76, normal exam. Labs pending.

**ASSESSMENT:** Stable hypertension with new symptoms.

**PLAN:** Adjust meds, follow-up in 1 week.

## 10.3 How to Judge Quality

- **Accuracy:** Compare output to input—ensure no added information or hallucinations (target: 95% fidelity).
- **Clarity:** Readability score (e.g., Flesch-Kincaid ↗ 60 for medical notes); use clinician feedback for coherence.
- **Evaluation Method:** Blind review by 3 physicians on 20 test cases, scoring 1-5 on accuracy/clarity; aim for average ↗4.

## 11 Ethics & Privacy

### 11.1 Offline Use & No Data Storage

The app runs entirely offline on local hardware, requiring no internet. By default, no patient data is stored—inputs are processed in memory and deleted after output generation. Clinicians can opt-in for temporary session storage, but it's encrypted and auto-deletes after 24 hours.

### 11.2 Risks & Limits

- **Risks:** Potential for misinterpretation if input is ambiguous; mitigated by clinician review.
- **Limits:** Not suitable for complex diagnostics; may struggle with rare medical terms without fine-tuning. Hardware requirements: Minimum 8GB RAM for 2B model.
- **Mitigations:** Built-in warnings for low-confidence outputs; regular model updates via open-source community.

## 12 Hackathon Story

### 12.1 Problem → Solution → Demo Flow

**Problem:** Clinicians drown in paperwork, cloud AI risks privacy. **Solution:** Local MedGemma app for fast, secure note generation. **Demo Flow:** (1) Show raw input entry, (2) Generate SOAP note live, (3) Edit and export demo, (4) Highlight offline/privacy features.

### 12.2 What Judges Should Remember

This is a practical, deployable tool that solves real clinician burnout with unbreakable privacy—built on open MedGemma, ready for any clinic without big tech dependencies.